

1. *Simple gradient calculation.* Consider a function,

$$J = z_1 e^{z_1 z_2}, \quad z_1 = a_1 w_1 w_2, \quad z_2 = a_2 w_1 + a_3 w_2^2,$$

- (a) Compute the partial derivatives,  $\partial J / \partial w_j$  for  $j = 1, 2$ .
- (b) Write pseudo-code for python function that, given  $\mathbf{w}$  and  $\mathbf{a}$  computes  $J(\mathbf{w})$  and  $\nabla J(\mathbf{w})$ .

2. *Gradient with a logarithmic loss.* Consider the loss function,

$$J(\mathbf{w}, b) := \sum_{i=1}^N (\log(y_i) - \log(\hat{y}_i))^2, \quad \hat{y}_i = \sum_{j=1}^p x_{ij} w_j + b,$$

This is an MSE loss function, but in log domain.

- (a) Find the gradient components,  $\partial J / \partial w_j$  and  $\partial J / \partial b$ .
- (b) Complete the following python function

```
def Jeval(w,b,...):
    ...
    return J, Jgradw, Jgradb
```

that computes  $J$  and  $\nabla_w J$  and  $\nabla_b J$ . You need to complete the arguments of the function. To receive full credit, avoid using for loops.

3. *Gradient with an inverse function.* Consider the nonlinear least squares fit loss function

$$J(\mathbf{w}) = \sum_{i=1}^n \left[ y_i - \frac{1}{w_0 + \sum_{j=1}^d w_j x_{ij}} \right]^2.$$

- (a) Compute the gradient components,  $\partial J / \partial w_j$ . You may want to define the intermediate variable,

$$z_i = w_0 + \sum_{j=1}^d w_j x_{ij}.$$

Also, you can write separate answers for  $\partial J / \partial w_0$  and  $\partial J / \partial w_j$  for  $j = 1, \dots, d$ .

- (b) Complete the following function to compute the loss and gradient,

```
def Jeval(w,...):
    ...
    return J, Jgrad
```

For the gradient, you may wish to use the function,

```
Jgrad = np.hstack((Jgrad0, Jgrad1))
```

to stack two vectors.

4. *Gradient with nonlinear parametrization.* Given data  $(x_i, y_i)$  with binary class labels  $y_i \in \{0, 1\}$ , consider the binary cross-entropy loss function,

$$J(\mathbf{a}, \mathbf{b}) := \sum_{i=1}^N \log(1 + e^{z_i}) - y_i z_i, \quad z_i = \sum_{j=1}^d a_j e^{-(x_i - b_j)^2 / 2}.$$

- (a) Compute the gradient components,  $\partial J/\partial a_j$  and  $\partial J/\partial b_j$ .
- (b) Complete the following function to compute the loss and gradient,

```
def Jeval(a,b,...):
    ...
    return J, Jgrada, Jgradb
```

Avoid for loops to receive full credit.

5. *Finding local and global minima.* Consider the function

$$f(x) = \frac{1}{4}x^2 + 1 - \cos(2\pi x).$$

- (a) Approximately draw  $f(x)$ .
  - (b) Write an equation for the gradient descent update to minimize  $f(x)$ .
  - (c) Using the graph in part (a), where is the global minima of  $f(x)$ ?
  - (d) Using the graph in part (a), find one initial condition where gradient descent could end up converging to a local minima that is not the global minima. The local minima do not have a closed form expression, but you should be able to use the graph in part (a) to “eyeball” an initial condition close to a bad local minima.
6. In this problem, we will see why gradient descent can often exhibit very slow convergence, even on apparently simple functions. Consider the objective function,

$$J(\mathbf{w}) = \frac{1}{2}b_1w_1^2 + \frac{1}{2}b_2w_2^2,$$

defined on a vector  $\mathbf{w} = (w_1, w_2)$  with constants  $b_2 > b_1 > 0$ .

- (a) What is the gradient  $\nabla J(\mathbf{w})$ ?
- (b) What is the minimum  $\mathbf{w}^* = \arg \min_{\mathbf{w}} J(\mathbf{w})$ ?
- (c) Part (b) shows that we can minimize  $J(\mathbf{w})$  easily by hand. But, suppose we tried to minimize it via gradient descent. Show that the gradient descent update of  $\mathbf{w}$  with a step-size  $\alpha$  has the form,

$$w_1^{k+1} = \rho_1 w_1^k, \quad w_2^{k+1} = \rho_2 w_2^k,$$

for some constants  $\rho_i$ ,  $i = 1, 2$ . Write  $\rho_i$  in terms of  $b_i$  and the step-size  $\alpha$ .

- (d) For what values  $\alpha$  will gradient descent converge to the minimum? That is, what step sizes guarantee that  $\mathbf{w}^k \rightarrow \mathbf{w}^*$ .
- (e) Take  $\alpha = 2/(b_1 + b_2)$ . It can be shown that this choice of  $\alpha$  results in the fastest convergence. You do not need to show this. But, show that with this selection of  $\alpha$ ,

$$\|\mathbf{w}^k\| = C^k \|\mathbf{w}^0\|, \quad C = \frac{\kappa - 1}{\kappa + 1}, \quad \kappa = \frac{b_2}{b_1}.$$

The term  $\kappa$  is called the *condition number*. The above calculation shows that when  $\kappa$  is very large,  $C \approx 1$  and the convergence of gradient descent is slow. In general, gradient descent performs poorly when the problems are ill-conditioned like this.

7. *Nested optimization.* Suppose we are given a loss function  $J(\mathbf{w}_1, \mathbf{w}_2)$  with two parameter vectors  $\mathbf{w}_1$  and  $\mathbf{w}_2$ . In some cases, it is easy to minimize over one of the sets of parameters, say  $\mathbf{w}_2$ , while holding the other parameter vector (say,  $\mathbf{w}_1$ ) constant. In this case, one could perform the following *nested* minimization: Define

$$J_1(\mathbf{w}_1) := \min_{\mathbf{w}_2} J(\mathbf{w}_1, \mathbf{w}_2), \quad \hat{\mathbf{w}}_2(\mathbf{w}_1) := \arg \min_{\mathbf{w}_2} J(\mathbf{w}_1, \mathbf{w}_2),$$

which represent the minimum and argument of the loss function over  $\mathbf{w}_2$  holding  $\mathbf{w}_1$  constant. Then,

$$\hat{\mathbf{w}}_1 = \arg \min_{\mathbf{w}_1} J_1(\mathbf{w}_1) = \arg \min_{\mathbf{w}_1} \min_{\mathbf{w}_2} J(\mathbf{w}_1, \mathbf{w}_2).$$

Hence, we can find the optimal  $\mathbf{w}_1$  by minimizing  $J_1(\mathbf{w}_1)$  instead of minimizing  $J(\mathbf{w}_1, \mathbf{w}_2)$  over  $\mathbf{w}_1$  and  $\mathbf{w}_2$ .

- (a) Show that the gradient of  $J_1(\mathbf{w}_1)$  is given by

$$\nabla_{\mathbf{w}_1} J_1(\mathbf{w}_1) = \nabla_{\mathbf{w}_1} J(\mathbf{w}_1, \mathbf{w}_2)|_{\mathbf{w}_2 = \hat{\mathbf{w}}_2}.$$

Thus, given  $\mathbf{w}_1$ , we can evaluate the gradient from (i) solve the minimization  $\hat{\mathbf{w}}_2 := \arg \min_{\mathbf{w}_2} J(\mathbf{w}_1, \mathbf{w}_2)$ ; and (ii) take the gradient  $\nabla_{\mathbf{w}_1} J(\mathbf{w}_1, \mathbf{w}_2)$  and evaluate at  $\mathbf{w}_2 = \hat{\mathbf{w}}_2$ .

- (b) Suppose we want to minimize a nonlinear least squares,

$$J(\mathbf{a}, \mathbf{b}) := \sum_{i=1}^n \left( y_i - \sum_{j=1}^d b_j e^{-a_j x_i} \right)^2,$$

over two parameters  $\mathbf{a}$  and  $\mathbf{b}$ . Given parameters  $\mathbf{a}$ , describe how we can minimize over  $\mathbf{b}$ . That is, how can we compute,

$$\hat{\mathbf{b}} := \arg \min_{\mathbf{b}} J(\mathbf{a}, \mathbf{b}).$$

- (c) In the above example, how would we compute the gradients,

$$\nabla_{\mathbf{a}} J(\mathbf{a}, \mathbf{b}).$$